

ANTHROPIC

Understanding AI Fluency

A Framework for Effective Human–AI Collaboration

“AI Fluency is the ability to work with AI systems in ways that are effective, efficient, ethical, and safe.”

Based on Anthropic's AI Fluency Course · Dakan & Feller, 2025 · May 2026

Course 1: <https://anthropic.skilljar.com/ai-fluency-framework-foundations>

Course 2: <https://anthropic.skilljar.com/ai-capabilities-and-limitations>

Reference: <https://ringling.libguides.com/ai/framework>

INTRODUCTION

What Is AI Fluency?

AI Fluency is the ability to work with AI systems in ways that are **effective, efficient, ethical, and safe**. It includes the practical skills, knowledge, insights, and values needed to adapt as AI technologies continue to evolve.

This framework focuses on **generative AI** — systems that produce new content (text, images, code, audio, video) rather than simply sorting or classifying existing content. The recommendation engine choosing your next video, the spam filter in your inbox: those are AI too, but that's not what this is about.

THE CORE IDEA

Fluent AI use is not about blind trust or wholesale scepticism. It is about holding an accurate mental model of how the machine works — clear enough that you can calibrate your trust task by task, output by output.

Three ways we work with AI

AUTOMATION

AI executes your instructions

AI performs specific tasks based on specific human instructions. You define what needs to be done; the AI executes it. Most prompts start here.

AUGMENTATION

You collaborate as thinking partners

Humans and AI work together iteratively, each contributing to the outcome. Back-and-forth dialogue. Both parties shape the result.

AGENCY

AI works independently on your behalf

You configure AI to operate independently, setting its knowledge and behaviour patterns rather than specifying each action. AI may interact with other systems or people.

The 4Ds framework applies across all three modes. Whichever mode you choose, you will need all four competencies.

The Four Core Competencies: The 4Ds

Every effective human–AI collaboration requires four competencies — the 4Ds. Each describes a different dimension of working well with AI. They are not sequential steps; they operate continuously and in combination.

D1 · DELEGATION

What should humans do vs AI?

Deciding what work is appropriate for you, for AI, or for you together — and how to distribute those tasks strategically. *Domain expertise is the foundation.*

D2 · DESCRIPTION

How do we communicate clearly?

Communicating with AI in ways that create a productive collaborative environment. Product, process, and performance. *Clear communication is the key lever.*

D3 · DISCERNMENT

How do we evaluate what AI gives us?

Thoughtfully and critically evaluating what AI produces, how it produces it, and how it behaves. *Critical evaluation is a non-negotiable responsibility.*

D4 · DILIGENCE

How do we ensure responsible use?

Using AI ethically, maintaining transparency, and taking accountability for AI-assisted work. *Transparency builds trust and integrity.*

HUMANS PROVIDE

What only you bring

Critical thinking, judgement, creativity, domain expertise, and ethical oversight. These are not being automated away — they become more important.

AI PROVIDES

What the machine brings

Speed, scale, pattern recognition, broad knowledge, and processing capacity. The combination of human and AI capabilities exceeds either alone.

The 4Ds are not built overnight. They evolve through practice, reflection, and continued learning. Apply them deliberately and they become instinct.

Delegation

Delegation is making thoughtful decisions about what work is appropriate for you to do, for AI to do, or for you and AI to do together — and how to distribute those tasks. It requires both domain expertise and an understanding of AI capabilities.

PROBLEM AWARENESS

What am I trying to do?

Clearly understanding your goals and the nature of the work before involving AI. Know the problem before you reach for the tool. Rushing to AI without this step produces the wrong output efficiently.

PLATFORM AWARENESS

What does this AI do well?

Understanding the capabilities and limitations of different AI systems. Not all AI tools are equal — knowing which to use for what task is itself a skill, and one worth cultivating explicitly.

TASK DELEGATION

Here's the plan.

Thoughtfully distributing work between humans and AI to leverage the strengths of each. Some tasks should stay human. Some suit AI. Many are best done in partnership.

“Effective delegation requires both domain expertise and an understanding of AI capabilities.”

KEY TAKEAWAY · DELEGATION

Domain expertise is the foundation. You cannot delegate well without knowing the territory. Leverage the combined strength of humans and AI — do not simply offload as much as possible to the machine.

Description

Description is the ability to communicate with AI systems in ways that create a productive collaborative environment. Clear communication up front saves time and consistently leads to better results.

PRODUCT DESCRIPTION

What do I want?

Defining what you want in terms of outputs, format, audience, and style. Be explicit about the end result. The AI cannot read your mind — it can only respond to what you actually say.

PROCESS DESCRIPTION

How should you approach this?

Defining how the AI works through your request — step-by-step instructions, reasoning approach, constraints to follow. Guide the process, not just the outcome.

PERFORMANCE DESCRIPTION

How should you behave?

Defining the AI's behaviour during collaboration: concise or detailed, challenging or supportive, formal or casual. Set the right mode for the task before you begin.

“Clear communication with AI systems up front saves time and leads to better results.”

KEY TAKEAWAY · DESCRIPTION

Context, examples, and conversation — not just prompts. Prompts are the input; the context you build around them is what produces consistently good outputs. Be specific. Give examples. Say what you want, not what you don't want.

Discernment

Discernment is the ability to thoughtfully and critically evaluate what AI produces, how it produces it, and how it behaves. You cannot accurately assess an AI output without some model of where that output came from.

PRODUCT DISCERNMENT

Is what it produced any good?

Evaluating the quality of what AI produces: accuracy, appropriateness, coherence, relevance. Does the output actually address the goal? Not every good-looking output is a correct one.

PROCESS DISCERNMENT

Did it get there the right way?

Evaluating how the AI arrived at its output — looking for logical errors, lapses in attention, or inappropriate reasoning steps. The path matters, not just the destination.

PERFORMANCE DISCERNMENT

Is it behaving as needed?

Evaluating how the AI behaves during your interaction: is its communication style effective? Is it being sycophantic, evasive, or appropriately direct?

“Discernment works hand-in-hand with Description in a continuous feedback loop.”

KEY TAKEAWAY · DISCERNMENT

Critical evaluation of AI output is a non-negotiable responsibility. Discernment skills mitigate AI limitations and enable capabilities. Until you understand precisely how an output was generated, do not blindly accept and deliver it.

Diligence

Diligence is using AI responsibly and ethically. It includes making thoughtful choices about AI systems, maintaining transparency about AI's role in your work, and taking full accountability for the outputs you use or share.

CREATION DILIGENCE

Am I using the right tool, right?

Being thoughtful about which AI systems you use and how you interact with them. The choice of platform, the data you share, and the tasks you assign are all part of responsible use.

TRANSPARENCY DILIGENCE

Is everyone who needs to know, informed?

Being honest about AI's role in your work with everyone who needs to know. Hiding AI involvement when it matters erodes trust — yours and the organisation's.

DEPLOYMENT DILIGENCE

Am I standing behind what I ship?

Taking responsibility for verifying and vouching for the outputs you use or share. When you put your name on AI-assisted work, you own it — fully.

“Being honest about AI's role in your work, checking what it gives you, standing behind what you ship: that's all part of AI fluency.”

KEY TAKEAWAY · DILIGENCE

Maintain accountability for AI interactions, outputs, and their impact. Transparency builds trust and integrity in your work. Remember: it is your responsibility to disclose how AI was used in producing any output or deliverable you share.

The Description–Discernment Loop

Of the four Ds, Description and Discernment are uniquely iterative. They do not happen once — they run in a continuous cycle until the output meets your standard.

As you work with AI in any of the three interaction modes — Automation, Augmentation, or Agency — you will move repeatedly between describing what you need and evaluating what comes back. Each round of discernment informs the next round of description.

STEP 1 · DESCRIPTION

Describe clearly

Define the product, process, and performance you need. Be explicit: what is the output? How should AI approach it? What behaviour do you expect during collaboration?

STEP 2 · DISCERNMENT

Evaluate critically

Evaluate the output for product quality (accuracy, coherence, relevance), reasoning path (logical errors, lapses), and AI behaviour (sycophancy, evasiveness, inappropriate tone).

STEP 3 · ITERATE

Refine and repeat

Adjust your description based on what you observed. Repeat the cycle until the output meets your standard. You — not the AI — decide when to exit the loop.

The loop runs until the output meets your standard — not until the AI says it's done. You decide when to exit the loop. You decide when the quality is sufficient to share.

A NOTE ON DILIGENCE WITHIN THE LOOP


You may decide early in the process which of the three interaction modes (Automation, Augmentation, Agency) to use. But regardless of mode, Diligence requires that you know and understand how the final output was generated before you put your name on it. The loop is also how you build that understanding.

How AI Works: The Machine Model

The 4Ds describe what you need to do. The machine properties explain why. Modern AI systems are not uniformly capable or unreliable — they are strong and weak along specific, predictable axes.

Understanding these properties lets you delegate more wisely, describe your intent more clearly, discern the quality of what comes back, and exercise the right level of diligence before you use it.

How an AI assistant is created



There are four core properties you navigate every time you use an AI. Each has a capability zone and an edge. The skill is learning where those edges are for your specific tasks.

The Four Core Properties of AI Systems

01

Next Token Prediction

Where do AI answers come from?

Capability zone

Excels at tasks resembling well-worn patterns: summarisation, reformatting, explaining common concepts, drafting structured text. When the path is familiar, the output is fluent and coherent.

Watch for

Hallucination — generating plausible-sounding text that is not accurate. The model is writing an answer word-by-word, not looking one up. Novelty, sparsity, and factual precision are the edges.

03

Working Memory

What is the AI attending to right now?

Capability zone

Everything in the context window is available: your instructions, uploaded documents, the conversation so far. The model works with your specific context when it fits.

Watch for

Hard limits — this is the property with a cliff, not a gradient. When the window fills, something falls off the edge. Between sessions, the window empties by default. Things work until they abruptly don't.

02

Knowledge

What does AI actually know?

Capability zone

Topics well-represented in training data: mainstream science, popular programming languages, widely-discussed history. Broad and deep on well-trodden ground.

Watch for

The knowledge cutoff. The model's world froze at the end of training. It doesn't know what happened after that date unless given tools. Rare, niche, contested, or recent topics are the edges.

04

Steerability

How much am I in control?

Capability zone

You can specify role, tone, format, length, rules — and the AI applies them. Short, concrete, verifiable instructions give the highest degree of control.

Watch for

Steerability $\neq$ understanding. The model follows instructions the same way it does everything else: by continuing a pattern. As instructions get longer, more abstract, or require logical precision, control degrades.

When Properties Interact

The four properties don't operate in isolation. Most real-world AI failures are two properties meeting. Understanding the combination tells you what went wrong and what to do differently.

Hallucinated citation

Next Token Prediction meets Knowledge

The model generates what looks plausible (token prediction) in an area where training data had a gap it doesn't know is there (knowledge). The citation sounds real. It isn't. Remedy: verify factual claims, especially citations and statistics, against primary sources.

Drift in a long conversation

Working Memory meets Steerability

Early context fades as the window fills (working memory), and later instructions overwrite earlier ones (steerability). The model seems to forget your original framing. Remedy: start fresh sessions; re-state key constraints explicitly.

Confidently wrong arithmetic

Next Token Prediction meets Steerability

The model produces fluent-looking numbers (token prediction) with no native sense of quantity — just patterns over digits (steerability). Remedy: ask the AI to use code or a tool for any non-trivial calculation, and verify outputs independently.

Agreeing with a flawed premise

Sycophancy meets Next Token Prediction

The trained disposition toward agreement (a fine-tuning artefact) combines with the tendency to continue your framing rather than challenge it (token prediction). Remedy: explicitly ask the AI to identify flaws or push back on your premise.

THE PRACTICAL UPSHOT

Fluent AI use isn't about memorising every failure mode. It is about holding a small model of the machine in your head — clear enough that when something goes wrong, you can recognise which kind of wrong it is, locate where on the continuum you've drifted, and respond accordingly.

Key Terminology

AI Fluency

The ability to work with AI systems in ways that are effective, efficient, ethical, and safe.

The 4Ds

Delegation, Description, Discernment, Diligence — the four core competencies of AI Fluency.

Generative AI

AI systems that create new content (text, images, code) rather than classifying existing data.

Large Language Model (LLM)

A generative AI system trained on vast text data to understand and generate human language.

Pre-training

The initial training phase where a model learns language patterns from enormous text corpora.

Fine-tuning

Additional training that shapes a pre-trained model into a helpful assistant.

Context Window

The fixed-size workspace holding everything the AI can currently attend to.

Knowledge Cutoff Date

The point after which the AI has no built-in knowledge, based on when training ended.

Hallucination

When AI confidently states something that sounds plausible but is incorrect.

Temperature

A setting controlling output randomness: higher = more creative, lower = more predictable.

Retrieval Augmented Generation (RAG)

A technique connecting AI to external knowledge sources to reduce hallucinations.

Prompt Engineering

The practice of designing effective prompts to produce desired outputs.

Chain-of-Thought Prompting

Encouraging AI to work through a problem step by step before answering.

Few-Shot Learning

Teaching AI by showing examples of the desired input–output pattern.

Sycophancy

A failure mode where AI agrees with or flatters the user rather than being accurate.

Bias

Systematic patterns in AI outputs that unfairly favour or disadvantage certain groups.

SUMMARY

Practical Takeaways

AI Fluency is a practised skill, not a one-time orientation. These principles distil the course into the habits that move the needle.

- 01 Before you prompt, think about Delegation.**
Know your goal. Know the AI's strengths. Choose the right tool for the right task — and keep some tasks human. Domain expertise is the non-negotiable foundation.
- 02 Your Description is your biggest lever.**
Be specific about what you want (product), how to approach it (process), and how to behave (performance). Vague prompts produce vague outputs. Context, examples, and conversation outperform bare prompts.
- 03 Discernment is always on.**
AI output is a draft, not a verdict. Scrutinise factual claims, citations, numbers, and anything recent. Verify before you share. Discernment skills mitigate AI limitations and enable capabilities.
- 04 Diligence is not optional.**
When AI helped, say so. When you ship AI-assisted work, you own it completely. Transparency builds trust and integrity in your work.
- 05 Manage the Description–Discernment loop actively.**
This is a multi-cycle iteration. Evaluate AI's output, approach, and behaviour — then adjust your description and repeat. You decide when to exit the loop, not the AI.
- 06 Context is the resource to protect.**
The context window is finite. It empties between sessions. Keep prompts and files lean. Start a fresh session when the AI loses the thread — the memory has filled, not the model deteriorated.
- 07 Build intentionally; practise continuously.**
The 4Ds evolve over time with deliberate practice. Build a personal AI policy, a fluency development plan, and a prompt library. Continue to learn, stay committed, and use what you discover.

*Source: Anthropic AI Fluency Framework and Foundations · Dakan & Feller, 2025 ·
anthropic.skilljar.com/ai-fluency-framework-foundations · Released under CC BY-NC-SA 4.0*